Module 4
IT3105

Simon Borøy-Johnsen
MTDT

October 20, 2015

1 Structure

The game board is represented as a two-dimensional array of Cell objects. A Cell object has a position, a value, and a boolean value telling whether the cell has been merged or not (this value is set to False before each player move). The cell value is given as the exponent over 2 in the cell ($1024 = 2^{10} \rightarrow value = 10$).

Making a move is done by iterating through all cells in the grid, staring from the edge the cells are moving towards. For each cell, its nearest non-empty neighbour is checked. If the value of the neighbour matches the cell, they are merged. If not, the cell moves to the farthest empty cell in the moving direction. By starting the iteration from where the cells are going, the game board can be changed on the go during the iteration.

2 Expectimax

I used expectimax with dynamic depth in my solution.

The depth is chosen before each player move;

```
def get_depth():
    if number_of_empty_cells > 3:
        return 1
    elif number_of_empty_cells > 1:
        return 2
    else:
        return 3
```

I found this choice of depth to be sufficient in order to achieve 2048. It might not be the best way of doing it, but it yields consistent results good enough to achieve the top score for this assignment.

My expectimax implementation consists of two main components (Python functions);

- evaluate_player_move(depth): This function returns the expected score of the current game board. For the four different directions to make a move, the outcome of making the move is evaluated. If the depth is 0, whatever move yields the highest score is returned. If not, the function evaluate_computer_move is ran (with same depth), and whatever move yields the highest expected score is returned.

- evaluate_computer_move(depth): In this function, all possible computer moves are simulated (a computer move is just spawning a 2 or 4, with probabilities 0.9 and 0.1, respectively, at a random empty cell). For every possible computer move, the evaluate_player_move (with one less depth) is ran. All the results of running evaluate_player_move are weighted, and the mean is returned.

3   Heuristic

I used a rather simple heuristic function in my implementation of the 2048 solver.

The board is organized into a "path", following the pattern in figure [1].

Figure 1: Path pattern

Each cell in the path is given a label, p_n, with p_0 being the upper left cell, p_1 the one below, and so on all the way to p_{15} in the upper right corner.

The score of the board is then calculated using this formula;

$$\sum_{n=0}^{15} value(p_n) * r^n \tag{1}$$

where $value(p_n)$ is the value at cell p_n, and r is a constant in the range $(0, 1)$ (set to 0.25 during the demo session).

This enforces a monotonic decreasing order along the path, as r^n decreases for each step. Example: With $r = 0.25$, a 512 in cell p_n would yield four times as much to the score as it would in cell p_{n+1}. A 512 in cell p_n would be equivalent to a 2048 in cell p_{n+1}. The AI will therefore do what it can to get the larger values as close to the start of the path as possible.